

Retrieval-Augmented Generation (RAG) combines information retrieval with language models.
Chunking splits documents into smaller pieces for embedding and search.
Cosine similarity measures how closely two embedding vectors align.
Fixed chunking uses character windows while recursive chunking respects sentence boundaries.
BGE and MiniLM are popular open-source embedding models for semantic search.